

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/221489022>

Spectral peak tracking and its use in speech recognition.

Conference Paper · January 2000

Source: DBLP

CITATIONS

17

READS

37

1 author:



[M. Padmanabhan](#)

Institute of Electrical and Electronics Engineers

83 PUBLICATIONS 1,438 CITATIONS

SEE PROFILE

SPECTRAL PEAK TRACKING AND ITS USE IN SPEECH RECOGNITION

M. Padmanabhan

IBM T. J. Watson Research Center P. O. Box 218,
Yorktown Heights, NY 10598
<http://www.research.ibm.com/people/m/mukund>

1. INTRODUCTION

In this paper, we address the issue of making use of spectral peak location information in a speech recognition system. The cepstral features that are used in most speech recognition systems, though perceptually motivated, do not explicitly model spectral peak trajectory information, which is a valuable clue to identifying the underlying phone. We present a study that examines the utility of using this information in speech recognition, to augment the information present in the cepstra.

We propose a method based on bandpass filtering the speech signal using several filters with different passbands, and using an adaptive IIR filter to track the locations of the spectral peaks in each bandpass output. This method has the advantage that along with the estimate of the spectral peak frequency, it also provides the energy at the spectral peaks (a feature that turns out to be quite informative). In initial experiments, the bandpass filters were chosen to correspond to the formant ranges, consequently, the locations of the spectral peaks are expected to correspond to the locations of the formants, for voiced sounds.

We next investigated the utility of using this spectral peak information to help discriminate between the phones used in speech recognition. In order to quantify the information provided by the new features (over and above the information provided by the cepstra), we measure the mutual information between the augmented feature vector (cepstra augmented with the new features) and the phonetic class labels, and compare it to the mutual information between the classes and the cepstra. Finally, we experimented with feature fusion techniques, where the new features were appended to the cepstra, and a new speech recognition system was trained on the augmented features.

2. ESTIMATION OF SPECTRAL PEAK FEATURES

We propose a method based on using an adaptive IIR filter for tracking the locations of the spectral peaks in the

speech signal. In order to simplify the task of the adaptive filter, we first isolate spectral regions of the input signal by passing it through a bank of bandpass filters, such that each region contains only one dominant spectral component. In initial experiments, we based the choice of bandpass filters on physio-acoustical studies that indicate that the spectral peaks correspond to formant frequencies, and further, the first three formant frequencies lie in the range 280-710 Hz, 870-2250 Hz, 2250-2890 Hz [1]. Consequently, the speech signal is first filtered using a bank of three linear phase bandpass filters, with the passbands corresponding to these formant ranges, and all filters having the same group delay. Subsequently, an adaptive IIR filter [2] is used to track the frequency at which the spectral energy is maximum within each passband.

As mentioned earlier, this method provides both the estimate of the spectral peaks as well as the energy at these peaks. For the voiced regions, this correlates roughly with the formant frequencies and energies, however, for the unvoiced regions, the adaptive filter essentially converges to the location of the spectral peaks in these regions. Further, as the adaptive filter is not allowed to change very abruptly, it also enforces a relatively smooth transition in the spectral peak locations over time.

2.1. Adaptive Filter

Denoting the outputs of the three bandpass filters as $y_1(t)$, $y_2(t)$ and $y_3(t)$, the adaptive filter stage identifies the spectral peak in the bandlimited spectra of $y_i(t)$. This adaptive filter structure has previously been described in [2]. The basic idea is as follows: the adaptive filter is a multiple notch IIR filter with the notch frequencies being directly related to the filter coefficients (each notch frequency depends on exactly one filter coefficient). This multiple notch transfer function is obtained by embedding several digital resonators in a feedback loop, with the notch frequencies corresponding to the resonator frequencies. The goal of the adaptive algorithm is to minimize the power of the output of the notch filter. For the case where there are N sinusoids in the input, and there are N notches in the filter transfer function, the

We would like to acknowledge the support of DARPA under Grant MDA972-97-C-0012 for funding part of this work.




Figure 1: Adaptive filter structure

optimal solution is when the notch frequencies are equal to the input sinusoidal frequencies.

For our case, we assume that each bandpass output, $y_i(t)$ has a single spectral peak and use an adaptive filter with a single notch to track each of the $y_i(t)$. This filter structure is shown in Fig 1. The transfer function from the input to various nodes in the filter structure are given below:

$$\begin{aligned} H_e(z) &= \frac{x_e}{x_{in}} = \frac{1 - (2 - k_i^2)z^{-1} + z^{-2}}{1 - (2 - k_i^2)(1 - G)z^{-1} + (1 - 2G)z^{-2}} \quad (1) \\ H_{fb,i}(z) &= \frac{x_{fb,i}}{x_{in}} = \frac{(2 - k_i^2)z^{-1} - 2z^{-2}}{1 - (2 - k_i^2)(1 - G)z^{-1} + (1 - 2G)z^{-2}} \quad (2) \\ H_{ps,i}(z) &= \frac{x_{ps,i}}{x_{in}} = \frac{k_i^2 z^{-1}}{1 - (2 - k_i^2)(1 - G)z^{-1} + (1 - 2G)z^{-2}} \quad (3) \end{aligned}$$

The transfer function $H_e(z)$ represents a notch filter, with the notch frequency ω_i being related to the filter coefficient, k_i , through the following equation

$$k_i = 2 \sin\left(\frac{\omega_i}{2}\right) \quad (4)$$

The algorithm for adapting the filter coefficient is given by

$$k_i(n+1) = k_i(n) - \mu \frac{x_e(n)x_{ps,i}(n)}{\langle x_{ps,i}(n)x_{ps,i}(n) \rangle + \epsilon} \quad (5)$$

The term $x_e(n)x_{ps,i}(n)$ denotes the pseudo-gradient of the objective function (the coefficient is adapted in a direction opposite the pseudo-gradient), and the $\langle x_{ps,i}(n)x_{ps,i}(n) \rangle + \epsilon$ denotes a power normalizing term that modifies the gradient direction to point in the Newton direction (i.e. the normalizing term approximates the inverse of the Hessian of the objective function). The spectral peak location can be inferred from the value of the filter coefficient k_i after it has converged using (4).

An additional feature of the filter structure is that the transfer function $H_{fb,i}$ is the complement of the notch transfer function, i.e., it represents a bandpass transfer function with the center frequency corresponding to the resonator frequency. Consequently, the power of the signal at $x_{fb,i}$ represents the power of the input speech signal at this frequency. Details of the comparison between the spectral peak estimates provided by the adaptive filter and formant estimates provided by Xwaves are provided in [4].

3. MUTUAL INFORMATION BETWEEN ACOUSTIC FEATURE AND PHONETIC CLASS

The usefulness of an acoustic feature may be measured by the amount of information it provides in discriminating between phonetic classes. This can be quantified by the mutual information between the feature vector and the phonetic class¹. Let c denote the phonetic class, and Z denote the acoustic feature vector. The mutual information between Z and c is defined by

$$I(Z; c) = \sum_c p(c) \int_Z p(Z/c) \log \left[\frac{p(Z/c)}{p(Z)} \right] dZ \quad (6)$$

Though (6) cannot be expressed in closed form, by vector quantizing Z and approximating the integral with a summation, it may be rewritten as

$$I(Z; c) = \sum_c p(c) \sum_{Z_j} p(Z_j/c) \log \left[\frac{p(Z_j/c)}{p(Z_j)} \right] \quad (7)$$

We are essentially interested in measuring the amount of information available by augmenting the usual cepstral feature vector with these new features. If Z represents the cepstral feature vector, and z represents one of the new features, then an augmented feature vector \hat{Z} can be formed by concatenating Z and z . By vector quantizing \hat{Z} into the same number of codewords as for Z , the mutual information between the augmented vector \hat{Z} and c , $I(\hat{Z}; c)$, can be computed from (8).

$$I(\hat{Z}; c) = \sum_c p(c) \sum_{\hat{Z}_j} p(\hat{Z}_j/c) \log \left[\frac{p(\hat{Z}_j/c)}{p(\hat{Z}_j)} \right] \quad (8)$$

The amount of information added by z can now be computed very simply

$$\delta I_z = I(\hat{Z}; c) - I(Z; c) \quad (9)$$

3.1. The amount of incremental information

The baseline acoustic feature, Z , was assumed to be the 13-dimensional Mel cepstral observation vector. We evaluated the δI_z associated with six new features, namely, the estimates of the first two formants from Xwaves [3], x_1, x_2 , the spectral peak locations of the first and second adaptive filters, s_1, s_2 and the energy at these frequencies, e_1, e_2 . The new features were observed either at the same time frame as the associated cepstra, or at any of the 4 preceding or following time frames, leading to 54 sets of features. The classes, c corresponded to 58 phonetic classes. The number of feature vectors per class was limited to 20000, with the

¹There also appears to be a link between the measure described here and the objective function for heteroschedastic discriminant analysis (HDA). This is explored further in [4].




Figure 2: Incremental mutual information between new features and phonetic classes

feature vectors for all classes totalling approximately 1 million. We K-means clustered the 1 million feature vectors to ≈ 2500 gaussians and used a likelihood metric to label each feature vector with the label of the closest gaussian. Subsequently, (8) was used to compute the mutual information between the class and the quantized feature vectors.

The unconditional class distribution had an entropy of 5.8 bits, and the mutual information associated with the cepstra alone, $I(Z; z)$ was 1.62 bits. Subsequently, we evaluated the incremental mutual information, δI_z , associated with each of the six new features at 9 time frames and plotted them in Fig 2.

Several observations may be made from Fig 2:

- (i) the most informative features are the the energy at the peak frequencies, e_1, e_2 . Further, the amount of additional information provided by this new feature is around 0.1 bits, which is not negligible when compared to the 1.62 bits of information that are provided by the cepstra.
- (ii) the spectral peak estimates, s_1, s_2 provide less information than the energy
- (iii) the formant estimates provided by Xwaves seem to add almost no information to the baseline cepstral feature. This is surprising because the formant estimates, x_1, x_2 , should nominally provide the same amount of information as the spectral peak features, s_1, s_2 . One possible reason for the discrepancy is that for the unvoiced segments of speech, the spectral peak locations tend to be influenced by the peak location in the preceding and following voiced segments and tend to vary smoothly between these values within the unvoiced segment, whereas the x_1, x_2 estimates vary more or less randomly for the unvoiced segments of speech. Consequently, even though the x_1, x_2 features do not provide information in disambiguating all phones, one could expect them to at least provide information to disambiguate speech regions with a well defined formant structure. To verify this




Figure 3: Incremental mutual information between new features and 8 vowels

conjecture, we computed the mutual information using (9), using data from only 8 voiced vowels, IY, IH, AE, AA, AH, UH and UW. The entropy of the prior distribution of these 8 classes is 3 bits, and the cepstra provide 1.03 bits of information. The incremental information provided by the new features is summarized in Fig 3. From this figure, it may be seen that the x_1, x_2 features do provide the maximum amount of information.

4. SPEECH RECOGNITION

Next we experimented with feature fusion techniques to incorporate the new features into a speech recognition system. The task that we experimented on is a Voicemail transcription task [5], which represents large vocabulary spontaneous telephone speech. The size of the vocabulary is 15k words, and the perplexity of the trigram LM that was used is approximately 100. The amount of acoustic training data available is approximately 70 hours. We present results on both a development test set comprising of 43 voicemail messages (approximately 20 mts of speech) and an evaluation test set comprising of 86 voicemail messages (approximately 40 mts of speech).

4.1. System overview

The speech recognition system uses a phonetic representation of the words in the vocabulary. Each phone is modelled with a 3-state left-to-right HMM. Further, we identify the variants of each state that are acoustically dissimilar by asking questions about the phonetic context in which the state occurs. The questions are arranged heirarchically in the form of a decision tree, and its leaves correspond to the basic acoustic units that we model. An acoustic feature vector is extracted every 10 ms, and we model the pdf of the




Figure 4: Word error rate vs feature dimension

feature vector for each leaf of the decision tree, with a mixture of gaussians (our system had 2313 leaves and 34k gaussians). We experimented with using a varying number of cepstral features, and augmenting the cepstral feature space with the new features, s_1 , s_2 , e_1 and e_2 . Subsequently, the final acoustic feature vector is obtained by augmenting the acoustic observation for a given frame with its first and second temporal derivatives.

4.2. Experimental Results

The word error rate results were computed on the Voicemail dev and eval test set [5] and are shown in Fig 4. The x-axis indicates the dimensionality of the extracted feature (either number of cepstra, or number of cepstra + e_1 , e_2 , or number of cepstra + s_1 , s_2). The figure shows that the (e_1, e_2) estimates do contain more information than the higher order (13th) cepstra and can be used to improve the performance of the system (by 5.7% on the dev test and 5.2% on the eval test). The observation from Fig 2 that the (e_1, e_2) are more informative than the frequency estimates (s_1, s_2) themselves is also supported by the recognition results. Also, all the experiments conducted so far have been on relatively clean speech, possibly for noisy speech, the roles of energy and frequency estimate will be reversed, but this still remains to be established.

5. DISCUSSION

In this paper, we experimented with augmenting the cepstral features used in a speech recognition system with spectral peak related features. The new features were estimated by passing the speech signal through a bank of bandpass filters, and using an adaptive filter to track the location of the spectral peak in each band. We also quantified the amount of

incremental information present in the new features by measuring the mutual information between the augmented feature and the phonetic classes and comparing it to the mutual information between the cepstra and the phonetic classes. Finally we incorporated the additional features into a speech recognition system, using feature fusion, and showed that the new features contained more information than the higher order cepstra and could help improve the word error rate.

There are however several issues which are still open. It is not clear that the bandpass frequency ranges (motivated by formant frequency ranges) that were used for these experiments are the most appropriate. Possible experiments for the future include subdividing the entire frequency range into uniform intervals and tracking the spectral peak within each frequency range. Also, the simple feature fusion approach that was explored in this paper could probably be considerably improved upon by the use of alternative classifier combination schemes [6, 7].

Acknowledgement

The author would like to thank Jun Huang for getting this work jump started during the summer of 1999, and Jing Huang and Karen Livescu for help with Xwaves.

6. REFERENCES

- [1] P. Ladefoged, "A Course in Phonetics", 1993, Harcourt Brace College Publishers, 301 Commerce Street, Suite 3700, Fort Worth, TX 76102.
- [2] M. Padmanabhan and K. Martin, "Resonator-based filter-banks for frequency domain applications", IEEE Trans. Circuits and Systems, Oct 1991.
- [3] Entropic Xwaves package.
- [4] M. Padmanabhan, "Use of spectral peak information in speech recognition", Proceedings of Speech Transcription Workshop, 2000, <http://www.itl.nist.gov/iaui/894.01/proc/>.
- [5] M. Padmanabhan et al., "Speech Recognition Performance on a VoiceMail Transcription Task", Proceedings of the Eurospeech 1999.
- [6] L. Yuan and D. Miller, "Ensemble Classification by Critic-driven Combining", Proceedings of the ICASSP 1999.
- [7] R. Schapire, Y. Freund, P. Bartlett, W. S. Lee, "Boosting the margin: A new explanation for the effectiveness of voting methods", Annals of Statistics, 26(5): 1651-1686, 1998.